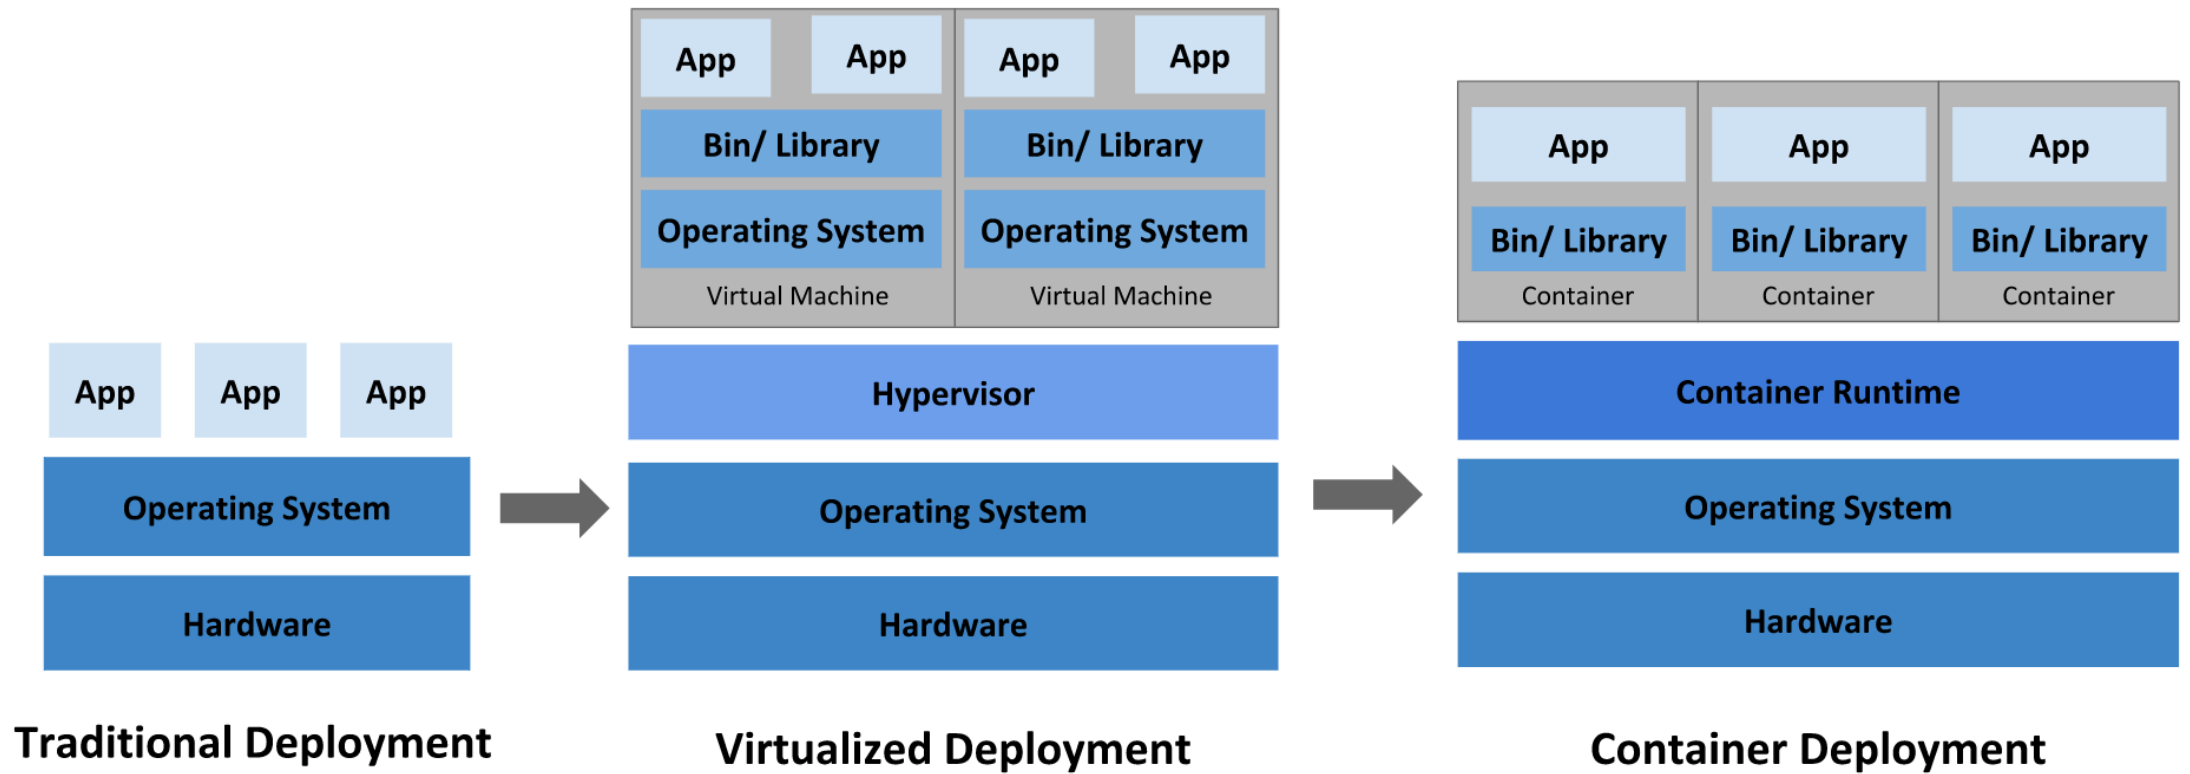



Introduction to Kubernetes

Part 2

Recap



Recap

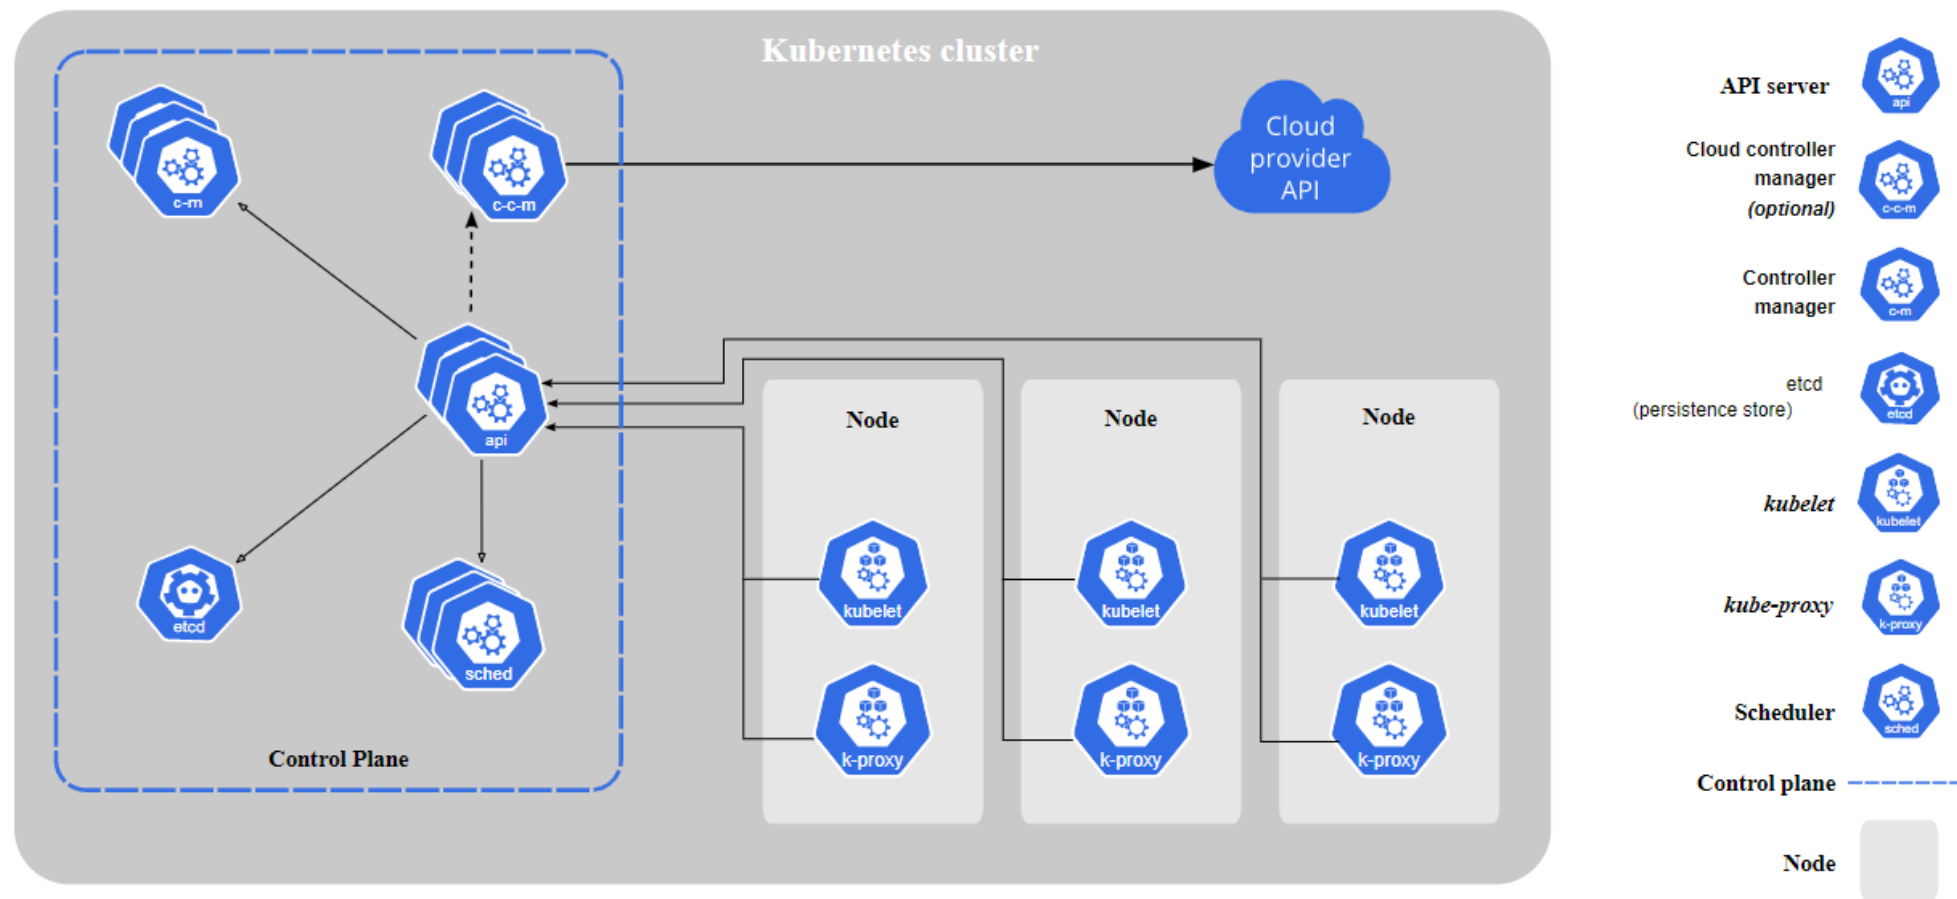


Recap

Kubernetes

- Distributed system to orchestrate containers
- **Declarative** and based on **control loops**
- Core object is **Pod**, which wraps 1..N containers
- **Namespace** – virtual cluster inside a physical one
- **Control** and **Data** planes

Recap



Course Plan

1. What is Kubernetes
2. What is Container
3. Kubernetes 10000-foot view
4. [Kubernetes core objects](#)
5. Daily interactions

Kubernetes Core Objects


- Pod
- Deployment
- StatefulSet
- Storage
- DaemonSet
- Job and CronJob
- Services and Ingress
- ConfigMaps and Secrets
- Policies

Kubernetes Core Objects

Pod:


- Smallest deployable unit
- Can have ≥ 1 containers
- Containers could be init-only
- All containers in a pod shares the same storage and network
- has unique IP

Kubernetes Core Objects




Kubernetes Core Objects


Sidecar



Adapter



Ambassador



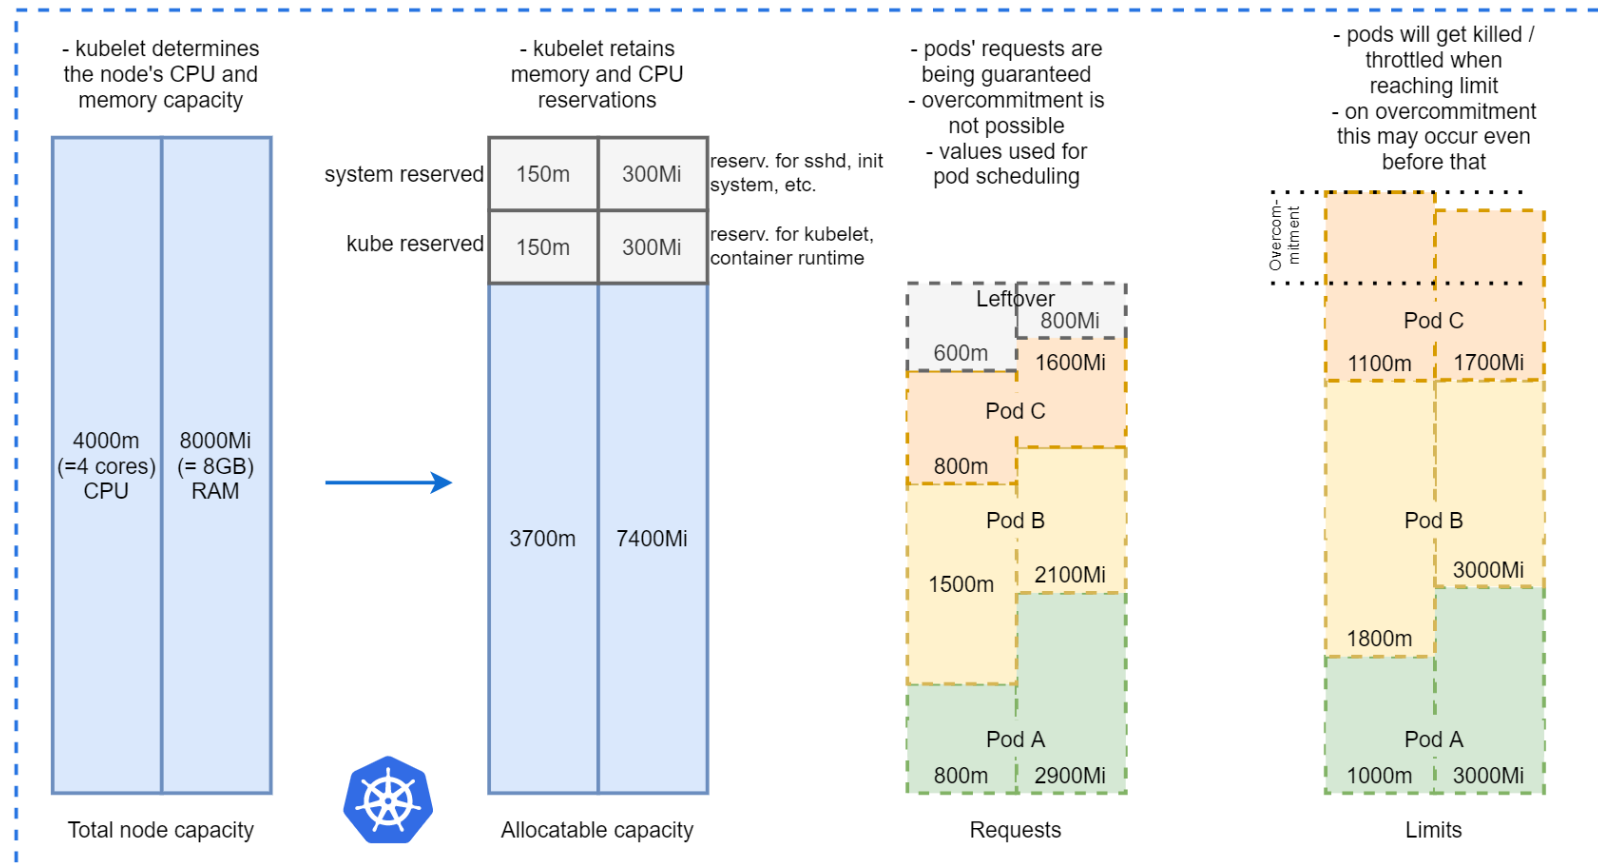
Kubernetes Core Objects

Resource-quotas: requests & limits

```
spec:
  containers:
  - name: nginx
    image: k8s.gcr.io/nginx-slim:0.8
    resources:
      limits:
        memory: 500Mi
        cpu: 200m
      requests:
        memory: 100Mi
        cpu: 100m
```

Kubernetes Core Objects

Kubernetes Resource Requests and Limits



Kubernetes Core Objects

To pull image images from Container Registry Kubernetes uses

- Node configuration
- or [imagePullSecret](#) property

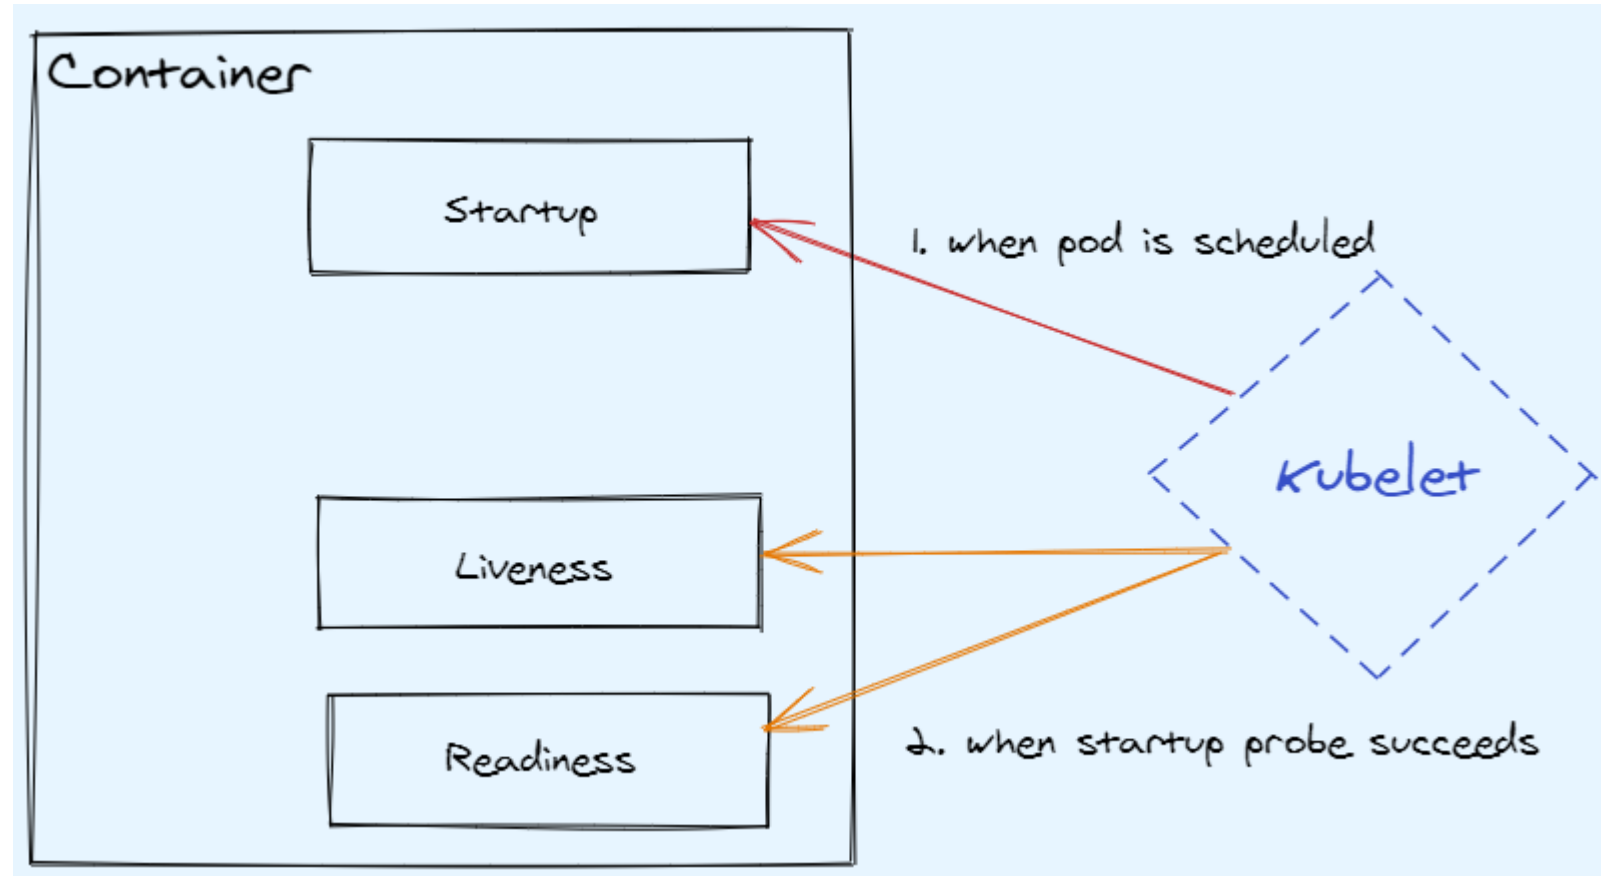
Kubernetes Core Objects

Pod Lifecycle:

- Get scheduled
- Remains on the node until completion, termination or deletion
- Do not self-heal (kubelet *heals*)
- Probes: startup, readiness, liveness

Kubernetes Core Objects

```
.....startupProbe:
.....httpGet:
.....  path: /
.....  port: 80
.....  initialDelaySeconds: 3
.....  periodSeconds: 3
.....livenessProbe:
.....httpGet:
.....  path: /
.....  port: 80
.....  initialDelaySeconds: 2
.....  periodSeconds: 3
.....readinessProbe:
.....httpGet:
.....  path: /
.....  port: 80
.....  initialDelaySeconds: 1
.....  periodSeconds: 3
```



Kubernetes Core Objects

Pod Scheduling:

- Available node resources (requests)
- Node selectors
- Node taints/toleration
- Affinity and anti-affinity

Kubernetes Core Objects

Deployment:

- a set of containers
- defines update-strategy
- scales container instances
- tools to watch the status and roll-back
- wraps [ReplicaSets](#)

Kubernetes Core Objects

ReplicaSet sets how many pods of exact specification Kubernetes should run (lower-level object)

Kubernetes Core Objects

Deployment Demo

Kubernetes Core Objects

StatefulSet: like Deployment, but with guarantees about ordering and naming (sticky identity)

Kubernetes Core Objects

Storage:

- **Volumes** - piece of storage available on the node, where pod is running. Volume is mounted into Pods
- **Persistent Volume (PV)** - a piece of storage in the cluster with reference to physical data location
- **Persistent Volume Claim (PVC)** - a storage request by a user. Often used to map StatefulSets to PVs
- **StorageClass** defines parameters of PV: provisioner (AzureDisk, AWSElasticBlockStore, etc), reclaim policy, allows resizing, etc

Kubernetes Core Objects

StatefulSet Demo

Kubernetes Core Objects

DaemonSet defines which Nodes should have an instance of given Pod

Kubernetes Core Objects

[Jobs/CronJob](#) run scheduled or one-time tasks

Kubernetes Core Objects

Service:

- not an application, but a record in etcd
- represents networking rules

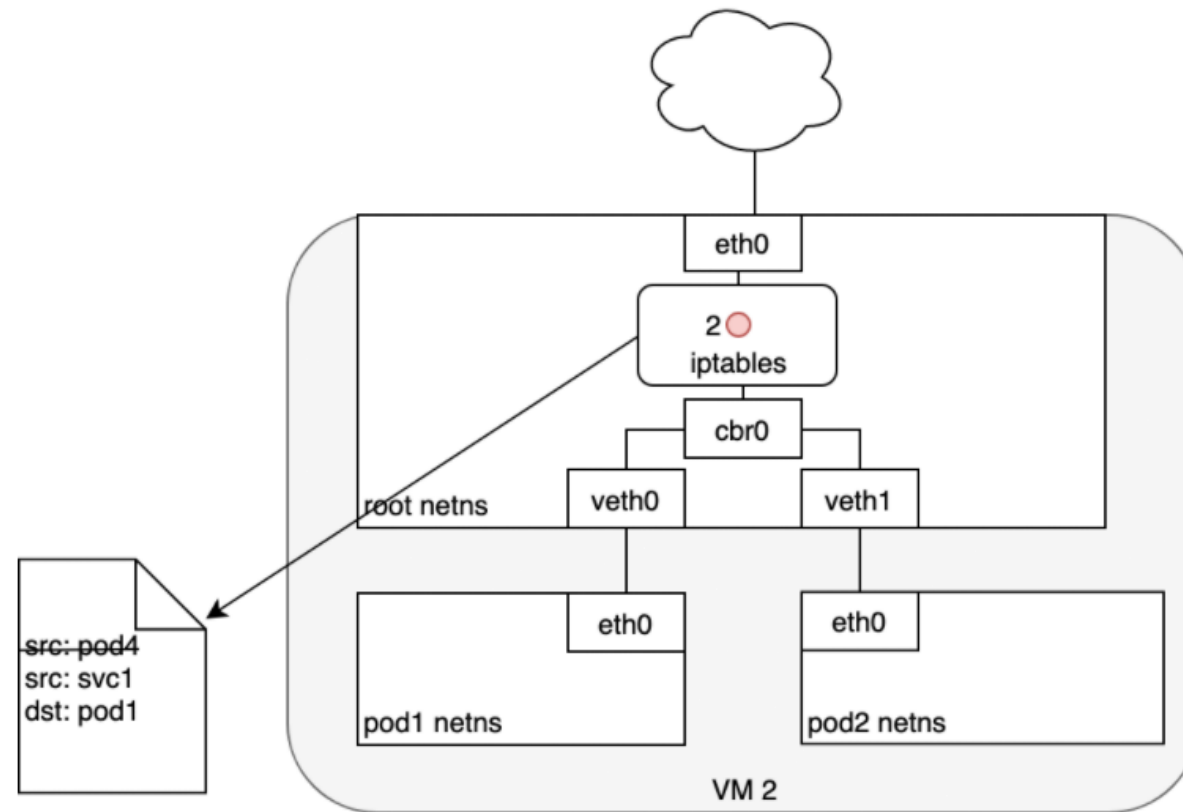
Kubernetes Core Objects

Service:

- pods are disposable, while a service is long-living
- service selects pods based on labels
- service might "route" traffic to pods
- [headless-service](#) - does not route traffic, but registers DNS identity

Kubernetes Core Objects

5.4 Life of a packet: Service to Pod



Kubernetes Core Objects

Service types:

- ClusterIP
- NodePort
- LoadBalancer

Kubernetes Core Objects


Services Demo

Kubernetes Core Objects

Ingress:

- API object (not a real application) to define rules
- Exposes HTTP/HTTPS routes from outside the cluster to [Services](#)
- Real work is done by [ingress-controller](#) (nginx, haproxy, envoy, etc)

Kubernetes Core Objects



Kubernetes Core Objects

Configuration:

- Decouple environment-specific information from application
- via environment variables
- via **ConfigMap** (CM) - API object to store **non-confidential data**. CM could be mapped as env-vars or volumes
- via **Secret** - API object to store **confidential data**
- Secrets are **not encrypted and are stored in etcd(!)**; secrets are encoded in base64, stored in tmpfs (RAM) instead of node-fs, could be encrypted at-rest and in-transit.

Kubernetes Core Objects

Policies:

- **LimitRanges**: min, max, default values for cpu/memory
- **Resource Quotas**: how many cpu, memory, storage could be used in a namespace
- **Pod Security Policy**: enforces pod configuration
- **NetworkPolicies**: enforces network boundaries (who can talk to whom); implemented by addons
- **RBAC**: defines identities and their permissions (who can do what)

Kubernetes Core Objects

- [Requests and Limits](#)
- [12-factor applications](#) and [Sidecar-containers patterns](#)
- [Stateful Applications in Kubernetes](#) article
- [Kubernetes Deconstructed](#) video
- [Understanding Kubernetes Networking](#) article
- Official [Pods](#), [Workloads](#), [Networking](#), [Storage](#), [Configuration](#), [Scheduling](#) documentation

Course Plan

1. What is Kubernetes
2. What is Container
3. Kubernetes 10000-foot view
4. Kubernetes core objects
5. Daily interactions

Daily Interactions

Navigate clusters and namespaces:

- ``kubectl config get-contexts``
- ``kubectl config set-context --current --namespace=kube-system``
- ``kubectl config use-context dev-cluster``
- ``kubectl get pods -n default``
- ``kubectl get pods -n default --context second``

If you have a lot of clusters and namespaces: [kubectx + kubens](#) tools

Daily Interactions

Basic operations:

- get help: ``kubectl --help`` and ``kubectl explain ...`` and ``kubectl api-resources``
- apply: ``kubectl apply -f sample_deployment.yaml``
- get & watch: ``kubectl get pods -o wide -w``
- describe: ``kubectl describe deployment nginx-deployment``
- scale: ``kubectl scale deployment nginx-deployment --replicas=1``
- delete: ``kubectl delete -f sample_deployment.yaml`` or ``kubectl delete deployment sample``

Daily Interactions

Troubleshooting:

- `get` or `describe` an object
- get `logs`: `kubectl logs pod_name container_name --tail=100`
- get `events`: `kubectl get events -n default`
- `port-forward` to pod/service: `kubectl port-forward svc/service1 port`
- run `shell in a container`: `kubectl exec pod-name -i -t -- bash`

If it is too simple, then welcome to [complete troubleshooting deployments](#) guide

Daily Interactions

Sharing, versioning, releasing Kubernetes Objects:

- Plain yaml files
- [kustomize](#)
- [helm](#)
- [Argo CD](#) and [Flux](#) (GitOps)

Daily Interactions

Additional resources:

- [Troubleshooting](#) VMware course
- [Kubernetes Best Practices](#) articles and videos by Google
- [Tasks](#) section at official docs

Alternative learning materials

- Official documentation is very well-written: <https://kubernetes.io/docs/concepts/>
- Full list of VMware courses: <https://kube.academy/courses>
- Azure learn-k8s materials: [azure-website](#) + [github-repo](#)
- [Kubernetes Up and Running](#) book. VMware [gives it in exchange for your personal data](#)
- [Designing Distributed Systems](#) book. Microsoft [gives it in exchange for your personal data](#)
- Learnk8s blog: <https://learnk8s.io/blog>